

# Docker for Machine Learning

- First steps – just use image
- More advanced – modify a public image
- More realistic – use (changed) images together

# Just use image

## Sources for Jupyter Notebooks

- <https://hub.docker.com/u/jupyter>
- <https://github.com/jupyter/docker-stacks>
- <https://github.com/Kaggle/docker-python>
- ...

Let's pick a universal one:

<https://hub.docker.com/r/jupyter/datascience-notebook>

# Just use image

## Steps

- Pull Image
- Run it

Can be one command

1) Download + Start

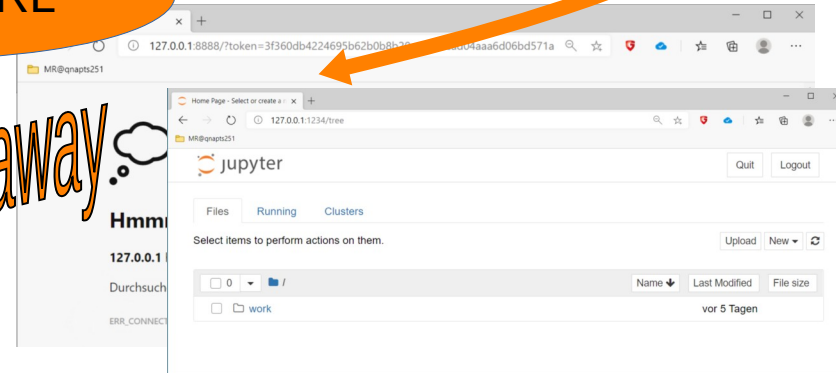
```
> docker run jupyter/datascience-notebook
```

3) Paste URL

2) Copy URL

*May or may not work right away*

```
PS E:\QSync\Mobile\Infos\Analytics\Big Data\Data Science\Machine Learning\ML DevOps> docker run jupyter/datascience-notebook
Unable to find image 'jupyter/datascience-notebook:latest' locally
latest: Pulling from jupyter/datascience-notebook
da7291352a9b: Pull complete
14428a6d4bcd: Pull complete
2c2d948710f2: Pull complete
ca78f2d1592f: Pull complete
2e7a2fbed339: Pull complete
d8b4c85d86bd: Pull complete
773273f9a979: Pull complete
73968a1cce87: Pull complete
b9bc7c3f9e37: Pull complete
3c47e33bf3e4: Pull complete
0f6d0494db95: Pull complete
62111c688dc9: Pull complete
1165e39c1e18: Pull complete
0e36996f6dd3: Pull complete
f9c15286a9d3: Pull complete
f79fd95f721: Pull complete
ec79857d12e8: Pull complete
556c3a8978ae: Pull complete
1696bccf21c7: Pull complete
729adc4b79ab: Pull complete
2c1bd9a56783: Pull complete
821bb95dab97: Pull complete
348dabe791df: Pull complete
2f8ef3706b64: Pull complete
eabb58054b1: Pull complete
38cb299b6e9d: Pull complete
d8528bd28762: Pull complete
Digest: sha256:a0cd0b721089799f9c8cc7b451a8a3d6ee6675691f9658a1b0ff856f9c0f9271b
Status: Downloaded newer image for jupyter/datascience-notebook:latest
Executing the command: jupyter notebook
[I 07:45:02.503 NotebookApp] Writing notebook server cookie secret to /home/jovyan/.local/share/jupyter/runtime/notebook_cookie_secret
[I 07:45:02.923 NotebookApp] JupyterLab extension loaded from /opt/conda/lib/python3.8/site-packages/jupyterlab
[I 07:45:02.923 NotebookApp] JupyterLab application directory is /opt/conda/share/jupyter/lab
[I 07:45:02.925 NotebookApp] Serving notebooks from local directory: /home/jovyan
[I 07:45:02.925 NotebookApp] Jupyter Notebook 6.1.6 is running at:
[I 07:45:02.925 NotebookApp] http://cc6c7ca6581b:8888/?token=3f360db4224695b62b0b8b20aea596cad04aaa6d06bd571a
```



# Just use image

## Troubleshooting

### 1. Stop running container

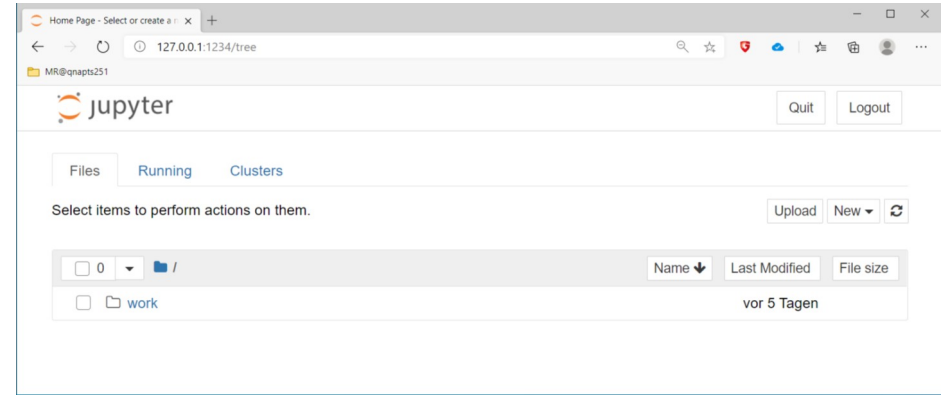
- `<Ctrl+C>`

### 1. Remove running container

- `docker container rm -f <mycontainer>`

### 2. Restart container with parameter(s)

- `docker run -p 1234:8888 jupyter/datascience-notebook`



# Just use image

## Working with Container

- Name container
  - `docker run -p 1234:8888 -name jupyter_ds --rm jupyter/datascience-notebook`
- Use data files (\*.csv, etc.)
  - `docker cp some.csv jupyter_ds:/home/jovyan/`
- Persist \*.ipynb, etc. outside container
  - `docker run -p 1234:8888 -v C:\Users\Mike_64\Documents:/home/jovyan --rm --name jupyter_ds jupyter/datascience-notebook`

# Modify a Public Image

Create individual Docker image with Dockerfile:

- Pull an image
- Customize it
- Save it to own Docker registry
- Run containers of the modified image

# Modify a Public Image

## Steps

- 1) Code Dockerfile
- 2) Build image
- 3) Run container off image

```
1 FROM jupyter/datascience-notebook
2
3
4 #base image is Ubuntu (https://github.com/docker/docker-stacks/issues/949)
5 #to install, user needs to be root
6 USER root
7
8 #update base packages
9 RUN apt-get update && \
10     apt-get --yes install \
11     apt-utils \
12     vim
13
14 #adding Python packages
15 RUN pip install mysql.connector sqlalchemy
16
17 #preparing Notebook environment
18 RUN mkdir /home/jovyan/.jupyter/
19 COPY jupyter_notebook_config.py /home/jovyan/.jupyter/
20
21 #back from root
22 USER $NB_UID
```

1

```
docker build -t mike_jn .
```

2

```
docker run -p 1234:1234 -v C:\Users\Mike_64\Documents:/home/jovyan --rm --name jupyter_ds mike_jn
```

```
[I 07:50:39.946 NotebookApp] JupyterLab application directory is /opt/conda/share/jupyter/lab
[I 07:50:39.948 NotebookApp] Serving notebooks from local directory: /home/jovyan
[I 07:50:39.948 NotebookApp] Jupyter Notebook 6.1.6 is running at:
[I 07:50:39.948 NotebookApp] http://e60e0fab498:1234/?token=b1e99e5af04b7d7d5f21c1ca8e94bd315abe47a8f3eee8e2
[I 07:50:39.948 NotebookApp] or http://127.0.0.1:1234/?token=b1e99e5af04b7d7d5f21c1ca8e94bd315abe47a8f3eee8e2
[I 07:50:39.948 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
```

3

```
To access the notebook, open this file in a browser:
file:///home/jovyan/.local/share/jupyter/runtime/nbserver-7-open.html
Or copy and paste one of these URLs:
http://e60e0fab498:1234/?token=b1e99e5af04b7d7d5f21c1ca8e94bd315abe47a8f3eee8e2
```

<http://127.0.0.1:1234/?token=b1e99e5af04b7d7d5f21c1ca8e94bd315abe47a8f3eee8e2>

# Modify a Public Image

Worth noting:

- Dockerfile name must be „Dockerfile“
- Docker build must be issued in directory of Dockerfile



# Modify a Public Image

For best practices:

[https://docs.docker.com/develop/develop-images/dockerfile\\_best-practices/](https://docs.docker.com/develop/develop-images/dockerfile_best-practices/)

A good tutorial:

<https://takacsmark.com/dockerfile-tutorial-by-example-dockerfile-best-practices-2018/>

# Use (Changed) Images Together

Pull / change multiple image to form „apps“:

- Pull images
- Customize them
- Save them to own Docker registry
- Run them as services with docker-compose

# Use (Changed) Images Together

# Steps

- 1) Code docker-compose.yml file
- 2) Run docker-compose

```

1 version: '3'
2 services:
3     mysql:
4         image: mike_ms
5         hostname: mike_ms
6         container_name: mike_ms
7         ports:
8             # <Port exposed> : < MySQL Port running inside container>
9             - '2345:3306'
10        expose:
11            # Opens port 3306 on the container
12            - '3306'
13            # Where our data will be persisted
14            #restart: always
15        environment:
16            MYSQL_DATABASE: 'dsdb'
17            # So you don't have to use root, but you can if you like
18            MYSQL_USER: 'dsdb_user'
19            # You can use whatever password you like
20            MYSQL_PASSWORD: 'dsdb_user_pwd'
21            # Password for root access
22            MYSQL_ROOT_PASSWORD: 'root_pwd'
23        volumes:
24            - C:\Users\Mike_64\Documents\MLOne_data:/var/lib/mysql

```

1

```
docker-compose up
```

2

# Use (Changed) Images Together

## Worth noting:

- docker-compose file name: „docker-compose.yml“ (\*.yaml)
- docker-compose up / down: issued in \*.yaml file directory
- Services may error out when started at once
  - start one service at a time:

```
docker-compose up mysql
```

```
docker-compose up jupyter
```

- `Docker-compose down` removes all containers
- Stop container(s) with  
`docker container mike_jn stop`

# Use (Changed) Images Together

## Working with services

- Name container
  - `docker run -p 1234:8888 -name jupyter_ds --rm jupyter/datascience-notebook`
- Use data files (\*.csv, etc.)
  - `docker cp some.csv jupyter_ds:/home/jovyan/`
- Persist \*.ipynb, etc. outside container
  - `docker run -p 1234:8888 -v C:\Users\Mike_64\Documents:/home/jovyan --rm --name jupyter_ds jupyter/datascience-notebook`